

## Machine\_Learning\_Classification\_Algorithms\_Role\_in\_Fetal\_Health\_Prediction

By : VENKATESH\_MUNGI, Email: [venkateshmungi@gmail.com](mailto:venkateshmungi@gmail.com) (<mailto:venkateshmungi@gmail.com>)

### Exploratory Data Analysis

```
In [1]: # importing the required libraries
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings("ignore")
```

```
C:\Users\VENKATESH MUNGI\AppData\Local\Programs\Python\Python310\lib\site-packages\numpy\_distributor_init.py:30: UserWarning: loaded more than 1 DLL from .libs:
C:\Users\VENKATESH MUNGI\AppData\Local\Programs\Python\Python310\lib\site-packages\numpy\.libs\libopenblas.EL2C6PLE4ZYW3ECEVIV30XXGRN2NRFM2.gfortran-win_amd64.dll
C:\Users\VENKATESH MUNGI\AppData\Local\Programs\Python\Python310\lib\site-packages\numpy\.libs\libopenblas.FB5AE2TYXYH2IJRDKGDGQ3XBKLT43H.gfortran-win_amd64.dll
  warnings.warn("loaded more than 1 DLL from .libs:")
```

```
In [2]: # Loading the dataset

df = pd.read_csv(r"C:\\PYTHON\\PANDAS\\main_projects\\fetal_health.csv")
df
```

```
Out[2]:
```

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations	prolongued_decelerations	abnormal_short_term_variability	mean_value_of_short_term_variability	percentage_of_time_with
0	120	0.000	0.000	0.000	0.000	0.0	0.0	73	0.5	
1	132	0.006	0.000	0.006	0.003	0.0	0.0	17	2.1	
2	133	0.003	0.000	0.008	0.003	0.0	0.0	16	2.1	
3	134	0.003	0.000	0.008	0.003	0.0	0.0	16	2.4	
4	132	0.007	0.000	0.008	0.000	0.0	0.0	16	2.4	
...	...	...	...	...	...	...	...	...	...	
2121	140	0.000	0.000	0.007	0.000	0.0	0.0	79	0.2	
2122	140	0.001	0.000	0.007	0.000	0.0	0.0	78	0.4	
2123	140	0.001	0.000	0.007	0.000	0.0	0.0	79	0.4	
2124	140	0.001	0.000	0.006	0.000	0.0	0.0	78	0.4	
2125	142	0.002	0.002	0.008	0.000	0.0	0.0	74	0.4	

2126 rows × 22 columns

In [3]: *# Shallow copying the dataset for our future convenience*

```
fetal = df.copy()
fetal
```

Out[3]:

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations	prolongued_decelerations	abnormal_short_term_variability	mean_value_of_short_term_variability	percentage_of_time_with
0	120	0.000	0.000	0.000	0.000	0.0	0.0	73	0.5	
1	132	0.006	0.000	0.006	0.003	0.0	0.0	17	2.1	
2	133	0.003	0.000	0.008	0.003	0.0	0.0	16	2.1	
3	134	0.003	0.000	0.008	0.003	0.0	0.0	16	2.4	
4	132	0.007	0.000	0.008	0.000	0.0	0.0	16	2.4	
...	...	...	...	...	...	...	...	...	...	
2121	140	0.000	0.000	0.007	0.000	0.0	0.0	79	0.2	
2122	140	0.001	0.000	0.007	0.000	0.0	0.0	78	0.4	
2123	140	0.001	0.000	0.007	0.000	0.0	0.0	79	0.4	
2124	140	0.001	0.000	0.006	0.000	0.0	0.0	78	0.4	
2125	142	0.002	0.002	0.008	0.000	0.0	0.0	74	0.4	

2126 rows × 22 columns

- Here i have done shallow copy of the dataset, because: if any changes made in the dataset that will not refelect in original dataset.

### ***Data Understanding***

- baseline value: Baseline Fetal Heart Rate (FHR) (beats per minute)
- accelerations: Number of accelerations per second
- fetal\_movement: Number of fetal movements per second
- uterine\_contractions: Number of uterine contractions per second
- light\_decelerations: Number of light decelerations (LDs) per second
- severe\_decelerations: Number of severe decelerations (SDs) per second
- prolonged\_decelerations: Number of prolonged decelerations (PDs) per second
- abnormal\_short\_term\_variability: Percentage of time with abnormal short term variability
- mean\_value\_of\_short\_term\_variability: Mean value of short term variability
- percentage\_of\_time\_with\_abnormal\_long\_term\_variability: Percentage of time with abnormal long term variability
- mean\_value\_of\_long\_term\_variability: Mean value of long term variability
- histogram\_width: Width of histogram made using all values from a record
- histogram\_min: Histogram minimum value
- histogram\_max: Histogram maximum value
- histogram\_number\_of\_peaks: Number of peaks in the exam histogram
- histogram\_number\_of\_zeroes: Number of zeros in the exam histogram
- histogram\_mode: Histogram mode

- histogram\_mean: Histogram mean
- histogram\_median: Histogram median
- histogram\_variance: Histogram variance
- histogram\_tendency: Histogram tendency
- fetal\_health: Encoded as 1-Normal; 2-Suspect; 3-Pathological. It is our very target column in the dataset.

In [4]: *# Metadata information*

```
fetal.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2126 entries, 0 to 2125
Data columns (total 22 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   baseline value                           2126 non-null   int64
1   accelerations                           2126 non-null   float64
2   fetal_movement                           2126 non-null   float64
3   uterine_contractions                     2126 non-null   float64
4   light_decelerations                     2126 non-null   float64
5   severe_decelerations                     2126 non-null   float64
6   prolonged_decelerations                  2126 non-null   float64
7   abnormal_short_term_variability          2126 non-null   int64
8   mean_value_of_short_term_variability     2126 non-null   float64
9   percentage_of_time_with_abnormal_long_term_variability 2126 non-null   int64
10  mean_value_of_long_term_variability       2126 non-null   float64
11  histogram_width                           2126 non-null   int64
12  histogram_min                             2126 non-null   int64
13  histogram_max                             2126 non-null   int64
14  histogram_number_of_peaks                 2126 non-null   int64
15  histogram_number_of_zeroes               2126 non-null   int64
16  histogram_mode                            2126 non-null   int64
17  histogram_mean                            2126 non-null   int64
18  histogram_median                         2126 non-null   int64
19  histogram_variance                        2126 non-null   int64
20  histogram_tendency                       2126 non-null   int64
21  fetal_health                             2126 non-null   int64
dtypes: float64(8), int64(14)
memory usage: 365.5 KB
```

- By metadata information it is clear that our data frame has eight continuous values and fourteen discrete values

In [5]: *# Statistics of the data*

```
fetal.describe().T.style.set_properties(**{'background-color': 'green', 'color': 'white', 'border-color': 'white'})
```

Out[5]:

	count	mean	std	min	25%	50%	75%	max
<b>baseline value</b>	2126.000000	133.303857	9.840844	106.000000	126.000000	133.000000	140.000000	160.000000
<b>accelerations</b>	2126.000000	0.003178	0.003866	0.000000	0.000000	0.002000	0.006000	0.019000
<b>fetal_movement</b>	2126.000000	0.009481	0.046666	0.000000	0.000000	0.000000	0.003000	0.481000
<b>uterine_contractions</b>	2126.000000	0.004366	0.002946	0.000000	0.002000	0.004000	0.007000	0.015000
<b>light_decelerations</b>	2126.000000	0.001889	0.002960	0.000000	0.000000	0.000000	0.003000	0.015000
<b>severe_decelerations</b>	2126.000000	0.000003	0.000057	0.000000	0.000000	0.000000	0.000000	0.001000
<b>prolongued_decelerations</b>	2126.000000	0.000159	0.000590	0.000000	0.000000	0.000000	0.000000	0.005000
<b>abnormal_short_term_variability</b>	2126.000000	46.990122	17.192814	12.000000	32.000000	49.000000	61.000000	87.000000
<b>mean_value_of_short_term_variability</b>	2126.000000	1.332785	0.883241	0.200000	0.700000	1.200000	1.700000	7.000000
<b>percentage_of_time_with_abnormal_long_term_variability</b>	2126.000000	9.846660	18.396880	0.000000	0.000000	0.000000	11.000000	91.000000
<b>mean_value_of_long_term_variability</b>	2126.000000	8.187629	5.628247	0.000000	4.600000	7.400000	10.800000	50.700000
<b>histogram_width</b>	2126.000000	70.445908	38.955693	3.000000	37.000000	67.500000	100.000000	180.000000
<b>histogram_min</b>	2126.000000	93.579492	29.560212	50.000000	67.000000	93.000000	120.000000	159.000000
<b>histogram_max</b>	2126.000000	164.025400	17.944183	122.000000	152.000000	162.000000	174.000000	238.000000
<b>histogram_number_of_peaks</b>	2126.000000	4.068203	2.949386	0.000000	2.000000	3.000000	6.000000	18.000000
<b>histogram_number_of_zeroes</b>	2126.000000	0.323612	0.706059	0.000000	0.000000	0.000000	0.000000	10.000000
<b>histogram_mode</b>	2126.000000	137.452023	16.381289	60.000000	129.000000	139.000000	148.000000	187.000000
<b>histogram_mean</b>	2126.000000	134.610536	15.593596	73.000000	125.000000	136.000000	145.000000	182.000000
<b>histogram_median</b>	2126.000000	138.090310	14.466589	77.000000	129.000000	139.000000	148.000000	186.000000
<b>histogram_variance</b>	2126.000000	18.808090	28.977636	0.000000	2.000000	7.000000	24.000000	269.000000
<b>histogram_tendency</b>	2126.000000	0.320320	0.610829	-1.000000	0.000000	0.000000	1.000000	1.000000
<b>fetal_health</b>	2126.000000	1.304327	0.614377	1.000000	1.000000	1.000000	1.000000	3.000000

```
In [6]: # Getting Median of the dataset
fetal.median()
```

```
Out[6]: baseline value          133.000
accelerations                0.002
fetal_movement               0.000
uterine_contractions         0.004
light_decelerations          0.000
severe_decelerations         0.000
prolongued_decelerations     0.000
abnormal_short_term_variability 49.000
mean_value_of_short_term_variability 1.200
percentage_of_time_with_abnormal_long_term_variability 0.000
mean_value_of_long_term_variability 7.400
histogram_width              67.500
histogram_min                93.000
histogram_max                162.000
histogram_number_of_peaks     3.000
histogram_number_of_zeroes    0.000
histogram_mode               139.000
histogram_mean               136.000
histogram_median             139.000
histogram_variance           7.000
histogram_tendency            0.000
fetal_health                  1.000
dtype: float64
```

```
In [7]: # Getting Mode of the dataset
fetal.mode()
```

```
Out[7]:
```

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations	prolongued_decelerations	abnormal_short_term_variability	mean_value_of_short_term_variability	percentage_of_time_with_ab
0	133	0.0	0.0	0.0	0.0	0.0	0.0	60	0.8	

1 rows × 22 columns

```
In [8]: # Range of the dataset
fetal_range = fetal.max() - fetal.min()

fetal_range
```

```
Out[8]: baseline value          54.000
accelerations                 0.019
fetal_movement                0.481
uterine_contractions          0.015
light_decelerations           0.015
severe_decelerations          0.001
prolongued_decelerations      0.005
abnormal_short_term_variability 75.000
mean_value_of_short_term_variability 6.800
percentage_of_time_with_abnormal_long_term_variability 91.000
mean_value_of_long_term_variability 50.700
histogram_width              177.000
histogram_min                109.000
histogram_max                116.000
histogram_number_of_peaks     18.000
histogram_number_of_zeroes    10.000
histogram_mode               127.000
histogram_mean               109.000
histogram_median             109.000
histogram_variance           269.000
histogram_tendency            2.000
fetal_health                  2.000
dtype: float64
```

```
In [9]: # getting variance of the dataset
fetal.var()
```

```
Out[9]: baseline value          9.684222e+01
accelerations                 1.494279e-05
fetal_movement                2.177701e-03
uterine_contractions          8.679323e-06
light_decelerations           8.762835e-06
severe_decelerations          3.283272e-09
prolongued_decelerations      3.480381e-07
abnormal_short_term_variability 2.955928e+02
mean_value_of_short_term_variability 7.801153e-01
percentage_of_time_with_abnormal_long_term_variability 3.384452e+02
mean_value_of_long_term_variability 3.167716e+01
histogram_width              1.517546e+03
histogram_min                8.738061e+02
histogram_max                3.219937e+02
histogram_number_of_peaks     8.698876e+00
histogram_number_of_zeroes    4.985198e-01
histogram_mode               2.683466e+02
histogram_mean               2.431602e+02
histogram_median             2.092822e+02
histogram_variance           8.397034e+02
histogram_tendency            3.731116e-01
fetal_health                  3.774589e-01
dtype: float64
```

```
In [10]: # Knowing Skewness and Kurtosis of the data  
from scipy.stats import skew  
from scipy.stats import kurtosis
```

```
In [11]: fetal.skew()
```

```
Out[11]: baseline value          0.020312  
accelerations          1.204392  
fetal_movement         7.811477  
uterine_contractions    0.159315  
light_decelerations     1.718437  
severe_decelerations    17.353457  
prolongued_decelerations 4.323965  
abnormal_short_term_variability -0.011829  
mean_value_of_short_term_variability 1.657339  
percentage_of_time_with_abnormal_long_term_variability 2.195075  
mean_value_of_long_term_variability 1.331998  
histogram_width        0.314235  
histogram_min          0.115784  
histogram_max          0.577862  
histogram_number_of_peaks 0.892886  
histogram_number_of_zeroes 3.920287  
histogram_mode         -0.995178  
histogram_mean         -0.651019  
histogram_median       -0.478414  
histogram_variance      3.219974  
histogram_tendency     -0.311632  
fetal_health           1.849934  
dtype: float64
```

```
In [12]: fetal.kurtosis()
```

```
Out[12]: baseline value          -0.292943
accelerations                    0.767648
fetal_movement                  64.260821
uterine_contractions            -0.635071
light_decelerations             2.517461
severe_decelerations            299.424142
prolongued_decelerations        20.515918
abnormal_short_term_variability -1.051030
mean_value_of_short_term_variability 4.700756
percentage_of_time_with_abnormal_long_term_variability 4.252998
mean_value_of_long_term_variability 4.131254
histogram_width                 -0.902287
histogram_min                   -1.290422
histogram_max                   0.632769
histogram_number_of_peaks       0.504211
histogram_number_of_zeroes      30.365084
histogram_mode                  3.009531
histogram_mean                  0.933427
histogram_median                0.667259
histogram_variance              15.131589
histogram_tendency              -0.652639
fetal_health                    2.091215
dtype: float64
```

```
In [13]: # Checkng for null values in the dataset
fetal.isnull().sum()
```

```
Out[13]: baseline value          0
accelerations                    0
fetal_movement                  0
uterine_contractions            0
light_decelerations             0
severe_decelerations            0
prolongued_decelerations        0
abnormal_short_term_variability 0
mean_value_of_short_term_variability 0
percentage_of_time_with_abnormal_long_term_variability 0
mean_value_of_long_term_variability 0
histogram_width                 0
histogram_min                   0
histogram_max                   0
histogram_number_of_peaks       0
histogram_number_of_zeroes      0
histogram_mode                  0
histogram_mean                  0
histogram_median                0
histogram_variance              0
histogram_tendency              0
fetal_health                    0
dtype: int64
```



- From the above informatio it is clear that features in this dataset doesn't has any null values

## Look into Important Features

```
In [14]: # Let's Look into Important Features in our dataset using SelectKBest
        """
        ANOVA f-test Feature Selection

        ANOVA is an acronym for “analysis of variance” and is a parametric statistical hypothesis test for determining whether the means from two or more samples of data (often three or more) come from the same distribution or not.

        An F-statistic, or F-test, is a class of statistical tests that calculate the ratio between variances values, such as the variance from two different samples or the explained and unexplained variance by a statistical test, like ANOVA. The ANOVA method is a type of F-statistic referred to here as an ANOVA f-test.

        Importantly, ANOVA is used when one variable is numeric and one is categorical, such as numerical input variables and a classification target variable in a classification task.

        The results of this test can be used for feature selection where those features that are independent of the target variable can be removed from the dataset.

        When the outcome is numeric, and [...] the predictor has more than two levels, the traditional ANOVA F-statistic can be calculated.

        The scikit-learn machine library provides an implementation of the ANOVA f-test in the f_classif() function. This function can be used in a feature selection strategy, such as selecting the top k most relevant features (largest values) via the SelectKBest class.
        """
```

```
Out[14]: '\nANOVA f-test Feature Selection\n\nANOVA is an acronym for “analysis of variance” and is a parametric statistical hypothesis test for determining whether the means from two or more samples of data (often three or more) come from the same distribution or not.\n\nAn F-statistic, or F-test, is a class of statistical tests that calculate the ratio between variances values, such as the variance from two different samples or the explained and unexplained variance by a statistical test, like ANOVA. The ANOVA method is a type of F-statistic referred to here as an ANOVA f-test.\n\nImportantly, ANOVA is used when one variable is numeric and one is categorical, such as numerical input variables and a classification target variable in a classification task.\n\nThe results of this test can be used for feature selection where those features that are independent of the target variable can be removed from the dataset.\n\nWhen the outcome is numeric, and [...] the predictor has more than two levels, the traditional ANOVA F-statistic can be calculated.\n\nThe scikit-learn machine library provides an implementation of the ANOVA f-test in the f_classif() function. This function can be used in a feature selection strategy, such as selecting the top k most relevant features (largest values) via the SelectKBest class.\n\n'
```

```
In [15]: from sklearn.feature_selection import SelectKBest, f_classif
```

```
In [16]: select = SelectKBest(score_func=f_classif, k=21) # K = Number of features to select
```

```
In [17]: fitting = select.fit(fetal.drop('fetal_health',axis=1),fetal['fetal_health'])
```

```
In [18]: x = pd.DataFrame(fitting.scores_)
```

```
In [19]: columns = pd.DataFrame(fetal.drop('fetal_health',axis=1).columns)
```

```
In [20]: fscores = pd.concat([columns,x],axis=1)
```

```
In [21]: fscores.columns = ['Attribute','Score']
```

```
In [22]: fscores.sort_values(by='Score',ascending=False)
```

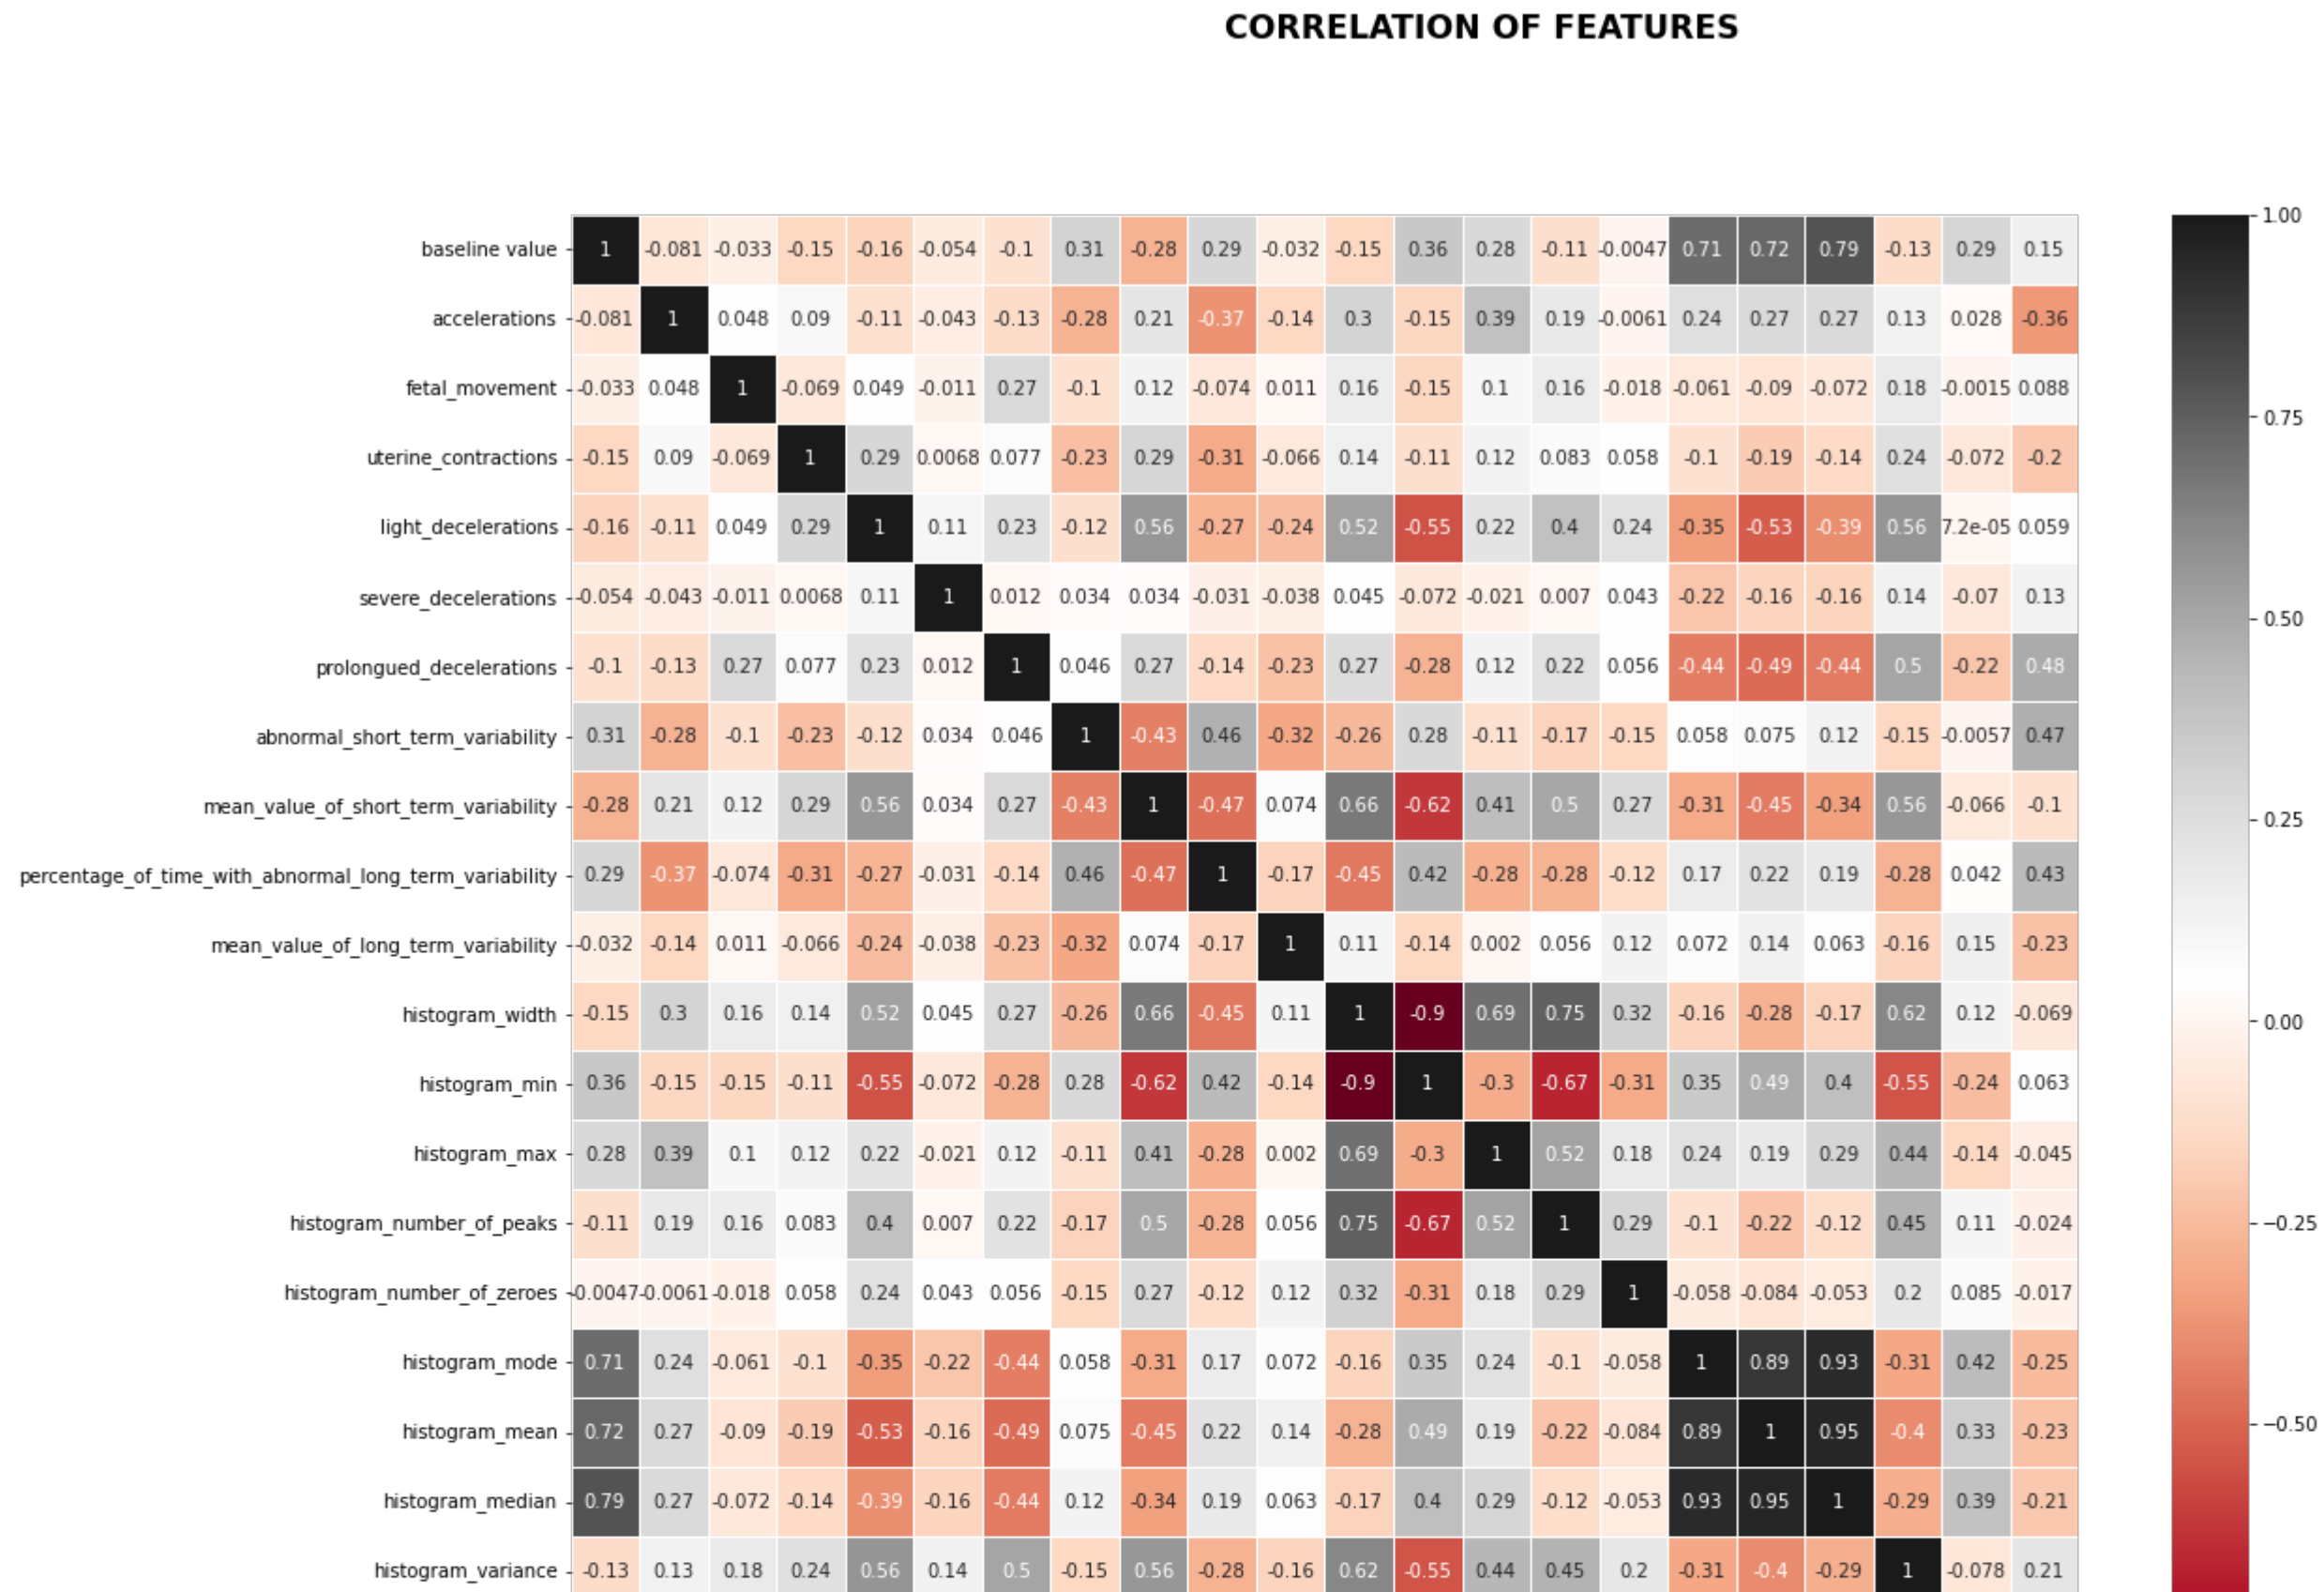
```
Out[22]:
```

	Attribute	Score
6	prolongued_decelerations	505.853206
9	percentage_of_time_with_abnormal_long_term_var...	345.156385
7	abnormal_short_term_variability	343.820419
17	histogram_mean	297.625497
16	histogram_mode	275.117696
18	histogram_median	248.772237
1	accelerations	196.027523
19	histogram_variance	150.796849
0	baseline value	140.621076
8	mean_value_of_short_term_variability	119.882006
3	uterine_contractions	93.715743
12	histogram_min	87.340503
10	mean_value_of_long_term_variability	70.174093
4	light_decelerations	66.864754
11	histogram_width	55.088241
20	histogram_tendency	44.542294
5	severe_decelerations	28.448156
14	histogram_number_of_peaks	12.104834
2	fetal_movement	11.679797
13	histogram_max	2.464923
15	histogram_number_of_zeroes	2.196373

- Fromtheaboveinformation it is clear that feature 'prolongued\_decelerations' has highest correlaation and 'histogram\_number\_of\_zeroes' has lowest correlation.
- First i want to use all features to build models. If the model accuracy is good its okay wtih all features, otherwise i would like to drop some features based on particular thresh-hold value.

**Let's Check Correlation of Variables with output variable "fetal\_health"**

```
In [23]: plt.figure(figsize = (18,15))
plt.suptitle("CORRELATION OF FEATURES", fontsize = 'xx-large', weight = 'extra bold' , ha = 'center')
corr = fetal.corr()
ax = sns.heatmap(corr, cmap='RdGy', annot=True, linewidths= 1.0)
plt.show()
```



- Observation:

histogram mean,median,mode have high correlations. The distribution may towards normal.

From the above correlation matrix, we can observe that the following features show some correlation with target variable fetal health:

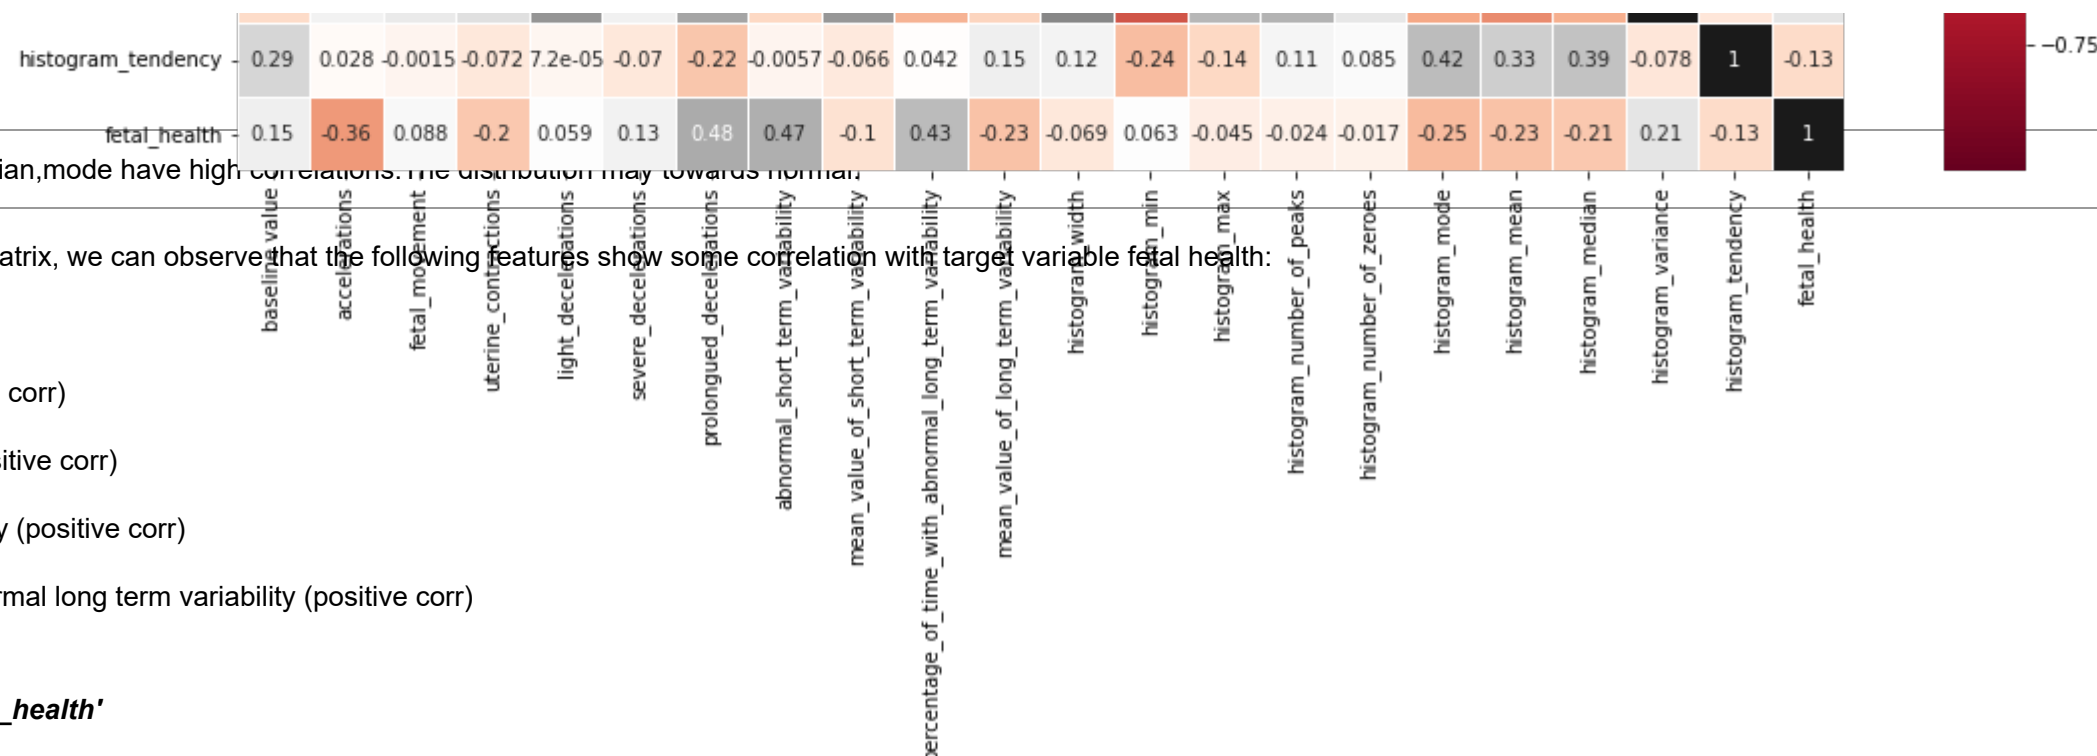
accelerations (negative corr)

uterine contractions (negative corr)

prolonged\_decelerations (positive corr)


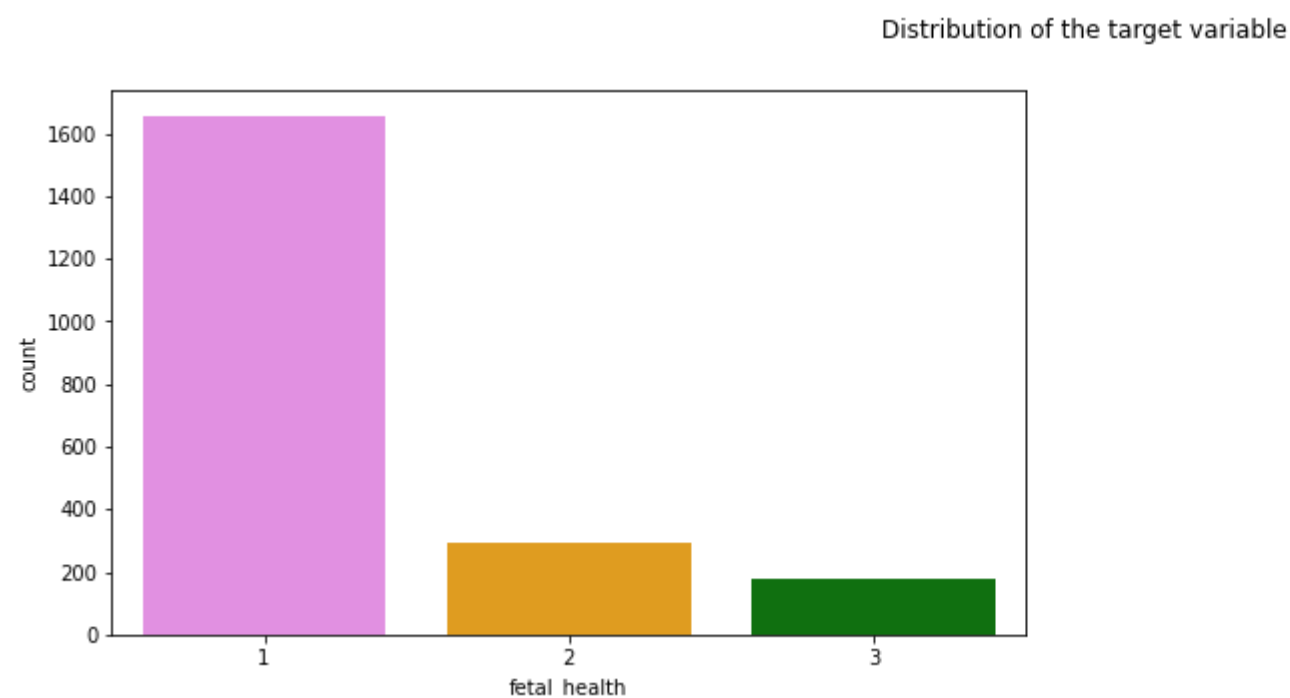
abnormal short term variability (positive corr)

percentage of time with abnormal long term variability (positive corr)



**Let's look into out put 'fetal\_health'**

```
In [24]: plt.figure(figsize=(18,5))
plt.subplot(1,2,1)
sns.countplot(x=fetal['fetal_health'], palette=['violet','orange','green'])
plt.subplot(1,2,2)
df['fetal_health'].value_counts().plot(kind='pie', autopct='%0.2f%%', explode=[0,0.1,0.1], startangle=90, colors=['violet','orange','green'],labels = ['Normal','Suspect','Pathological'])
plt.suptitle('Distribution of the target variable')
plt.show()
```



**Observation:**

- As class 3 "Pathological" has very little frequency it means that we are facing a “rare diseases” dataset. Also, this means that we can not remove any class and we have to study all of them.

- An highly Imbalanced Dataset. Which is obvious as Normal would be dominant.

**Solution:**

- Oversampling of the minority classes. To make better predictions.



```
In [25]: grouped = fetal.groupby(by='fetal_health').mean()
grouped
```

```
Out[25]:
```

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations	prolongued_decelerations	abnormal_short_term_variability	mean_value_of_short_term_variability	percentage_of_
<b>fetal_health</b>										
1	131.981873	0.003992	0.007963	0.004781	0.001941	6.042296e-07	0.000051	42.465861	1.430634	
2	141.684746	0.000275	0.008332	0.002390	0.000536	0.000000e+00	0.000095	61.901695	0.638983	
3	131.687500	0.000392	0.025676	0.003784	0.003670	3.409091e-05	0.001273	64.539773	1.575568	

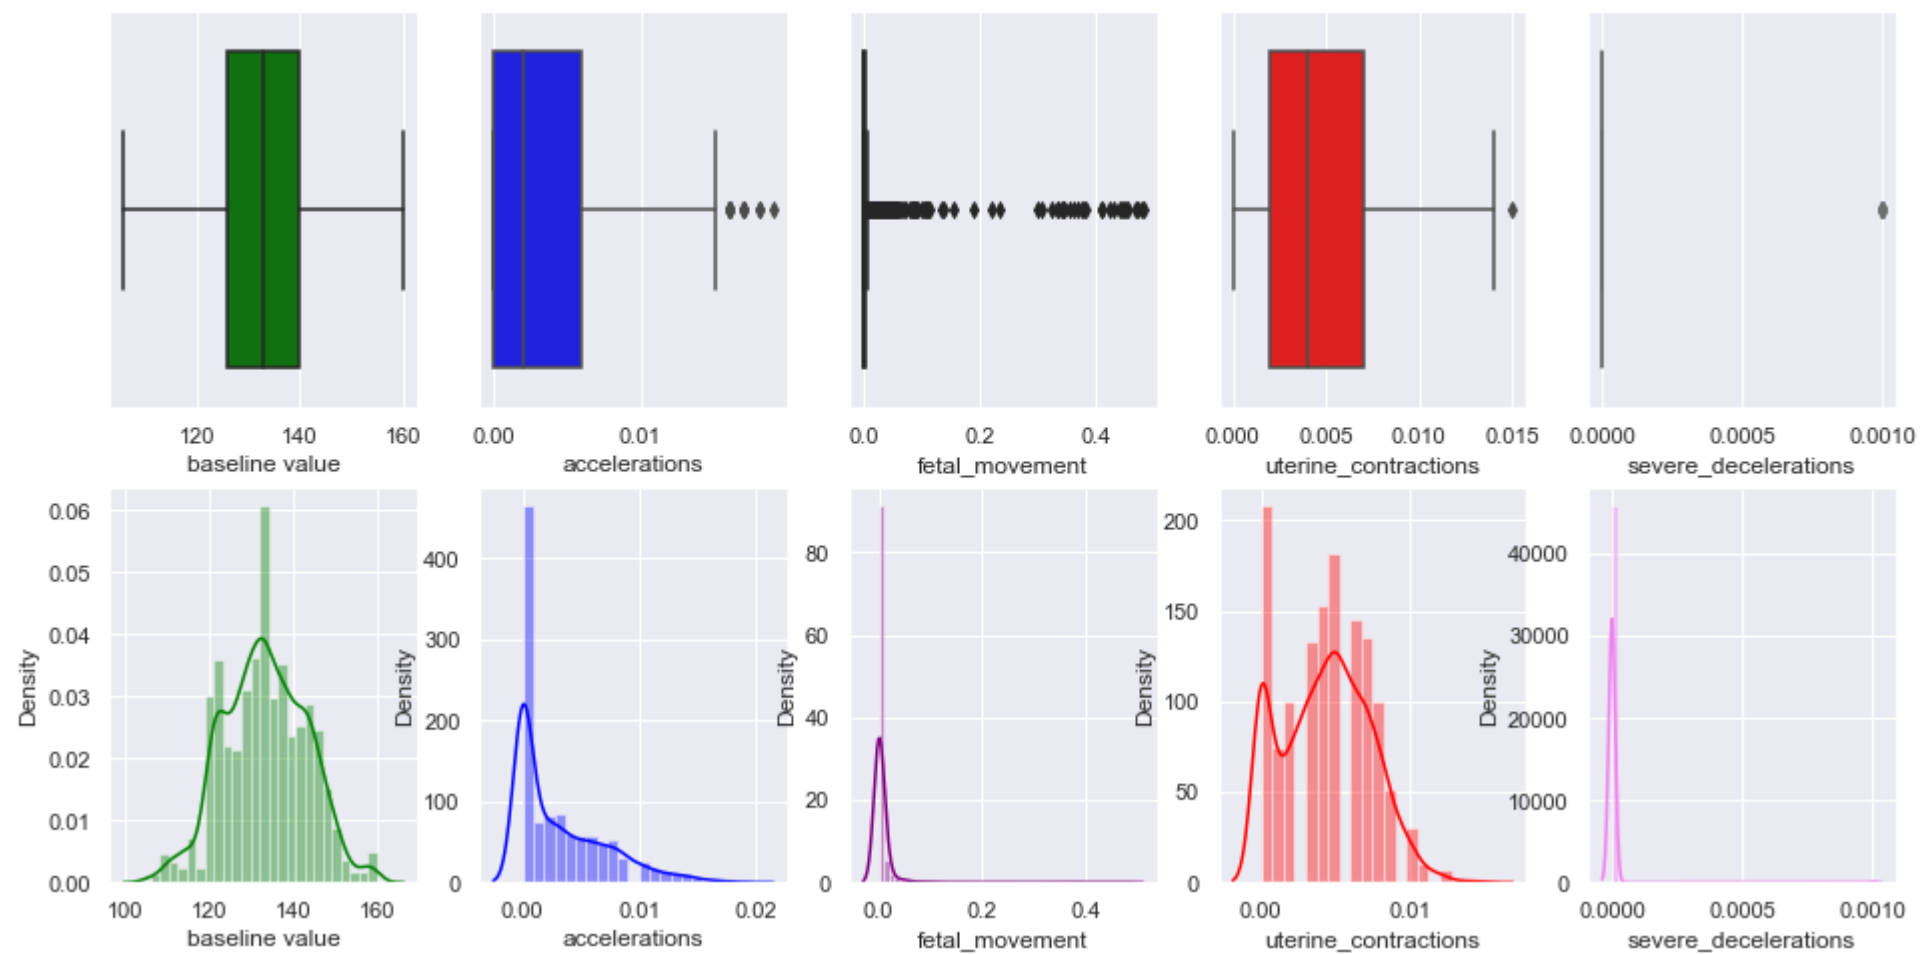
3 rows × 21 columns

```
In [26]: import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
sns.set()
for index,i in enumerate(grouped.columns,start=1):
    plt.figure(figsize=(6,4))
    sns.barplot(data=grouped,x=grouped.index,y=grouped[i])
    plt.show()
```

**Checking for outliers in our data**

```
In [27]: plt.figure(figsize = (16,8))
plt.suptitle("Checking for Outliers and Distribution pattern of variables", fontsize = 'x-large',weight = 'extra bold',ha = "center")
plt.subplot(2,5,1)
sns.boxplot(fetal['baseline value'], color = 'green')
plt.subplot(2,5,2)
sns.boxplot(fetal['accelerations'], color = 'blue')
plt.subplot(2,5,3)
sns.boxplot(fetal['fetal_movement'], color = 'purple')
plt.subplot(2,5,4)
sns.boxplot(fetal['uterine_contractions'], color = 'red')
plt.subplot(2,5,5)
sns.boxplot(fetal['severe_decelerations'], color = 'violet')
plt.subplot(2,5,6)
sns.distplot(fetal['baseline value'], color = 'green')
plt.subplot(2,5,7)
sns.distplot(fetal['accelerations'], color = 'blue')
plt.subplot(2,5,8)
sns.distplot(fetal['fetal_movement'], color = 'purple')
plt.subplot(2,5,9)
sns.distplot(fetal['uterine_contractions'], color = 'red')
plt.subplot(2,5,10)
sns.distplot(fetal['severe_decelerations'], color = 'violet')
plt.show()
```

**Checking for Outliers and Distribution pattern of variables**




- From the above graph it is clear that all features has outliers except 'baseline value.'

- All features are right skewed except 'baseline value'

```
In [28]: plt.figure(figsize = (16,8))
plt.suptitle("Checking for Outliers and Distribution pattern of variables", fontsize = 'x-large',weight = 'extra bold',ha = "center")
plt.subplot(2,5,1)
sns.boxplot(fetal['prolongued_decelerations'], color = 'green')
plt.subplot(2,5,2)
sns.boxplot(fetal['abnormal_short_term_variability'], color = 'blue')
plt.subplot(2,5,3)
sns.boxplot(fetal['mean_value_of_short_term_variability'], color = 'purple')
plt.subplot(2,5,4)
sns.boxplot(fetal['percentage_of_time_with_abnormal_long_term_variability'], color = 'red')
plt.subplot(2,5,5)
sns.boxplot(fetal['mean_value_of_long_term_variability'], color = 'violet')
plt.subplot(2,5,6)
sns.distplot(fetal['prolongued_decelerations'], color = 'green')
plt.subplot(2,5,7)
sns.distplot(fetal['abnormal_short_term_variability'], color = 'blue')
plt.subplot(2,5,8)
sns.distplot(fetal['mean_value_of_short_term_variability'], color = 'purple')
plt.subplot(2,5,9)
sns.distplot(fetal['percentage_of_time_with_abnormal_long_term_variability'], color = 'red')
plt.subplot(2,5,10)
sns.distplot(fetal['mean_value_of_long_term_variability'], color = 'violet')
plt.show()
```

### Checking for Outliers and Distribution pattern of variables

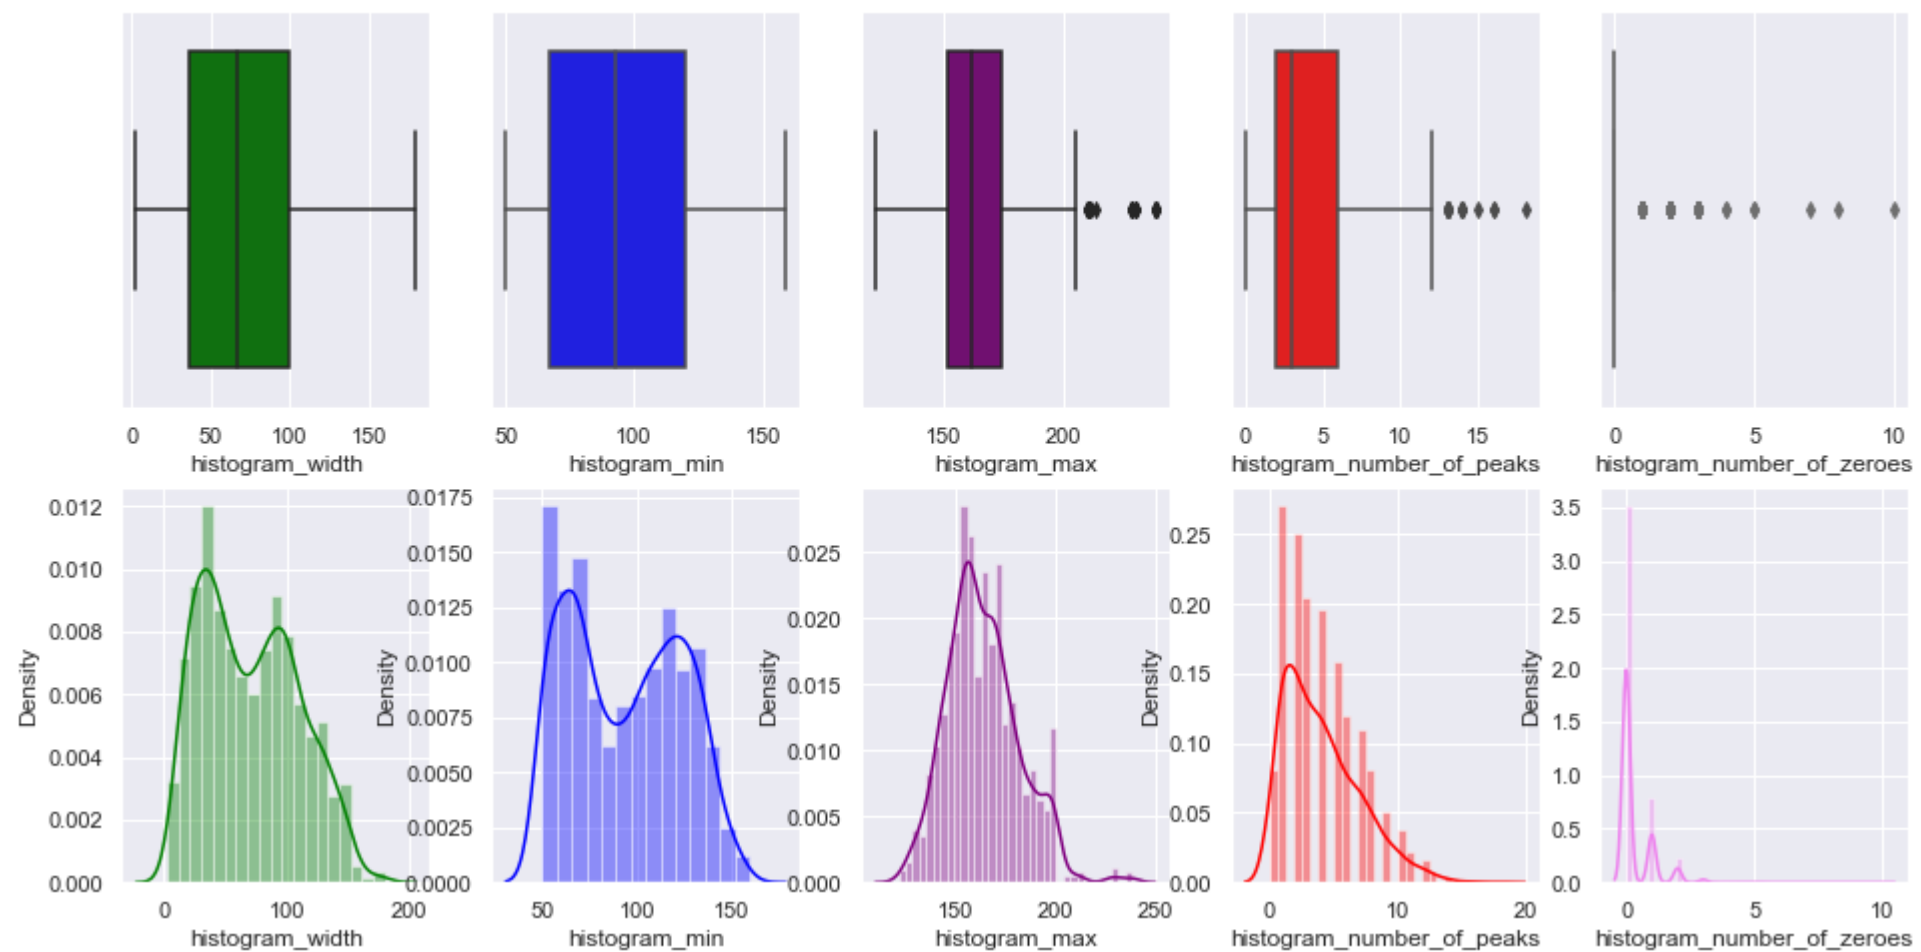




- From the above picture it is clear that all the features are right skewed and have outliers except 'abnormal\_short\_term\_variability'

```
In [29]: plt.figure(figsize = (16,8))
plt.suptitle("Checking for Outliers and Distribution pattern of variables", fontsize = 'x-large',weight = 'extra bold',ha = "center")
plt.subplot(2,5,1)
sns.boxplot(fetal['histogram_width'], color = 'green')
plt.subplot(2,5,2)
sns.boxplot(fetal['histogram_min'], color = 'blue')
plt.subplot(2,5,3)
sns.boxplot(fetal['histogram_max'], color = 'purple')
plt.subplot(2,5,4)
sns.boxplot(fetal['histogram_number_of_peaks'], color = 'red')
plt.subplot(2,5,5)
sns.boxplot(fetal['histogram_number_of_zeroes'], color = 'violet')
plt.subplot(2,5,6)
sns.distplot(fetal['histogram_width'], color = 'green')
plt.subplot(2,5,7)
sns.distplot(fetal['histogram_min'], color = 'blue')
plt.subplot(2,5,8)
sns.distplot(fetal['histogram_max'], color = 'purple')
plt.subplot(2,5,9)
sns.distplot(fetal['histogram_number_of_peaks'], color = 'red')
plt.subplot(2,5,10)
sns.distplot(fetal['histogram_number_of_zeroes'], color = 'violet')
plt.show()
```

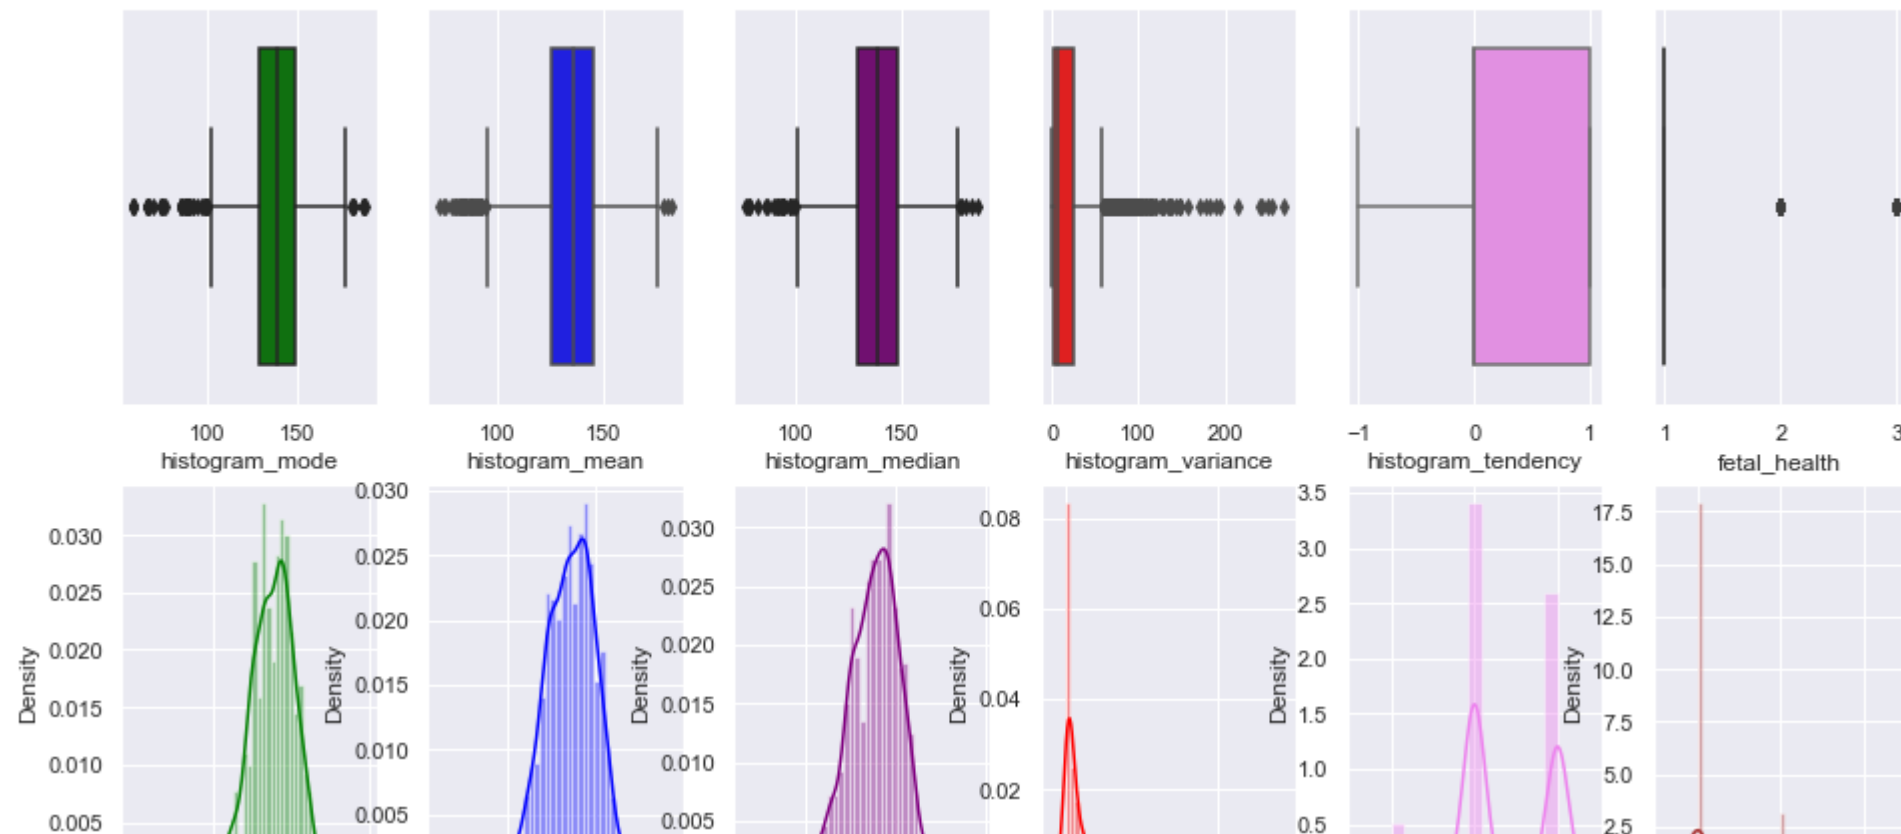
### Checking for Outliers and Distribution pattern of variables




- From the above visual it is clear that the features 'histogram width', 'histogram min' has no outliers remaining features has outliers and right skewness.

```
In [30]: plt.figure(figsize = (16,8))
plt.suptitle("Checking for Outliers and Distribution pattern of variables", fontsize = 'x-large',weight = 'extra bold',ha = "center")
plt.subplot(2,6,1)
sns.boxplot(fetal['histogram_mode'], color = 'green')
plt.subplot(2,6,2)
sns.boxplot(fetal['histogram_mean'], color = 'blue')
plt.subplot(2,6,3)
sns.boxplot(fetal['histogram_median'], color = 'purple')
plt.subplot(2,6,4)
sns.boxplot(fetal['histogram_variance'], color = 'red')
plt.subplot(2,6,5)
sns.boxplot(fetal['histogram_tendency'], color = 'violet')
plt.subplot(2,6,6)
sns.boxplot(fetal['fetal_health'], color = 'brown')
plt.subplot(2,6,7)
sns.distplot(fetal['histogram_mode'], color = 'green')
plt.subplot(2,6,8)
sns.distplot(fetal['histogram_mean'], color = 'blue')
plt.subplot(2,6,9)
sns.distplot(fetal['histogram_median'], color = 'purple')
plt.subplot(2,6,10)
sns.distplot(fetal['histogram_variance'], color = 'red')
plt.subplot(2,6,11)
sns.distplot(fetal['histogram_tendency'], color = 'violet')
plt.subplot(2,6,12)
sns.distplot(fetal['fetal_health'], color = 'brown')
plt.show()
```

### Checking for Outliers and Distribution pattern of variables





Even though , outliers capping is important , I don't want to capping outliers, because each and evry value in the health data set is import tant to analyse patient's health.

If we change the value to new value, then there is no meaning to that particular record.

In this project, I would like to continue without outlier capping of data

## Minority Class Oversampling

```
In [31]: x = fetal.iloc[:,0:21]
y = fetal.iloc[:, -1]
```

```
In [32]: from imblearn.over_sampling import RandomOverSampler
oversample = RandomOverSampler(sampling_strategy='not majority')
x_over, y_over = oversample.fit_resample(x, y)
```


```
In [33]: x_over.shape
```

```
Out[33]: (4965, 21)
```

```
In [34]: y_over.shape
```

```
Out[34]: (4965,)
```

```
In [35]: plt.figure(figsize=(18,6))
plt.suptitle('Data Before Over-sampling and after Over-Sampling')
plt.subplot(1,2,1)
sns.countplot(x=fetal['fetal_health'], palette=['#845ec2', '#ec4646', '#00af91'])
plt.subplot(1,2,2)
sns.countplot(data=pd.DataFrame(y_over),x='fetal_health', palette=['#845ec2', '#ec4646', '#00af91'] )
plt.show()
```



- From the above picture we can see that all the classes in our output feature are balanced.

## Feature Scalling

```
In [36]: from sklearn.preprocessing import StandardScaler
sdd = StandardScaler()
scaled_fetal = sdd.fit_transform(x_over)
```

```
In [37]: scaled_fetal
```

```
Out[37]: array([[ -1.52261655, -0.52042455, -0.23379762, ..., -0.77138339,
         1.13918796,  1.14402939],
        [ -0.32021627,  1.4995686 , -0.23379762, ...,  0.26415691,
        -0.30603768, -0.3865839 ],
        [ -0.22001624,  0.48957202, -0.23379762, ...,  0.15515266,
        -0.28234545, -0.3865839 ],
        ...,
        [ -0.22001624, -0.52042455, -0.23379762, ..., -1.75242157,
         1.66041688, -0.3865839 ],
        [  1.28298411, -0.52042455, -0.15181638, ...,  0.80917812,
        -0.59034436,  1.14402939],
        [  1.48338415, -0.52042455, -0.16821263, ...,  1.08168872,
        -0.56665214, -0.3865839 ]])
```

```
In [38]: new_fetal = pd.DataFrame(scaled_fetal, columns = x_over.columns)
```

```
In [39]: new_fetal
```

```
Out[39]:
```

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations	prolongued_decelerations	abnormal_short_term_variability	mean_value_of_short_term_variability	percentage_of_time_wit
0	-1.522617	-0.520425	-0.233798	-1.131941	-0.615206	-0.096703	-0.466968	0.982911	-0.707679	
1	-0.320216	1.499569	-0.233798	0.731792	0.315003	-0.096703	-0.466968	-2.294205	0.895893	
2	-0.220016	0.489572	-0.233798	1.353037	0.315003	-0.096703	-0.466968	-2.352725	0.895893	
3	-0.119816	0.489572	-0.233798	1.353037	0.315003	-0.096703	-0.466968	-2.352725	1.196563	
4	-0.320216	1.836234	-0.233798	1.353037	-0.615206	-0.096703	-0.466968	-2.352725	1.196563	
...	...	...	...	...	...	...	...	...	...	
4960	-1.522617	-0.520425	-0.233798	-0.821319	1.555282	-0.096703	1.509584	0.456232	-0.106340	
4961	-0.019616	-0.520425	-0.020646	-0.821319	-0.615206	-0.096703	-0.466968	0.748831	-0.908126	
4962	-0.220016	-0.520425	-0.233798	1.974281	0.315003	-0.096703	1.509584	0.280672	1.597456	
4963	1.282984	-0.520425	-0.151816	-0.510697	-0.615206	-0.096703	-0.466968	0.924391	-0.908126	
4964	1.483384	-0.520425	-0.168213	-1.131941	-0.615206	-0.096703	-0.466968	1.509590	-0.807902	

4965 rows × 21 columns

```
In [40]: import pickle
import requests
import json
import joblib
```

```
In [41]: joblib.dump(sdd, "sdd.joblib") # Dumping standard Scaling for model deployment
```

```
Out[41]: ['sdd.joblib']
```

## Train\_Test\_Split

```
In [42]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x_over, y_over, train_size = 0.70, shuffle = True, stratify = y_over)
```

## Model Building

### *Logistic Regression*

```
In [43]: from sklearn.linear_model import LogisticRegression
```

```
In [44]: model1 = LogisticRegression(multi_class='multinomial', solver = 'newton-cg', random_state = 0)
model1.fit(x_train, y_train)
```

```
Out[44]: LogisticRegression
LogisticRegression(multi_class='multinomial', random_state=0,
                  solver='newton-cg')
```

```
In [45]: lr_y_pred = model1.predict(x_test)
```

```
In [46]: from sklearn.metrics import confusion_matrix
cm1 = confusion_matrix(y_test, lr_y_pred)           # Confusion Matrix
cm1.ravel()
```

```
Out[46]: array([394,  88,  14,  50, 384,  63,   5,  46, 446], dtype=int64)
```

```
In [47]: from sklearn.metrics import log_loss

pred1 = model1.predict_proba(x_test)
logloss1 = log_loss(y_test, pred1)                 # Log_Loss Score
```

```
In [48]: from sklearn.metrics import accuracy_score           # Accuracy_Score

acc1 = accuracy_score(y_test, lr_y_pred)
```

- When LogisticRegression(solver = 'lbfgs', random\_state = 0) the accuray score is 75%
- When LogisticRegression(multi\_class='multinomial', solver = 'lbfgs', random\_state = 0) the accuray score is 77%
- When LogisticRegression(multi\_class='multinomial', solver = 'newton-cg', random\_state = 0) the accuray score is 83%
- When LogisticRegression(multi\_class='ovr', solver = 'newton-cg', random\_state = 0) the accuray score is 81%

Among all parameter tuning multi\_class='multinomial', solver = 'newton-cg' gives best accuracy score. So, I choosen third one for this model.



```
In [49]: from sklearn import metrics                                     # Precision_Score

precision_positive1 = metrics.precision_score(y_test, lr_y_pred, pos_label = 'positive', average='micro')

precision_negative1 = metrics.precision_score(y_test, lr_y_pred, pos_label = 'negative', average='micro')

In [50]: recall_sensitivity1 = metrics.recall_score(y_test, lr_y_pred, pos_label = 'positive', average='micro') # Recall_Score

recall_specificity1 = metrics.recall_score(y_test, lr_y_pred, pos_label = 'negative', average='micro')

In [51]: f1_positive1 = metrics.f1_score(y_test, lr_y_pred, pos_label = 'positive', average='micro')          # f1_Score

f1_negative1 = metrics.f1_score(y_test, lr_y_pred, pos_label = 'negative', average='micro')
```

- Use micro-averaging score when there is a need to weight each instance or prediction equally.
- Use macro-averaging score when all classes need to be treated equally to evaluate the overall performance of the classifier with regard to the most frequent class labels.

```
In [52]: #roc auc score
from sklearn.metrics import roc_curve, roc_auc_score
rocscore1 = roc_auc_score(y_test, pred1, multi_class='ovr', average= 'weighted')
```

- For Plotting ROC curves for all models is time consuming and will take too much space. Sinc, I would like to plot multiclass auc curve for best model later on.

#### **Note:**

OvR and OvO strategies can (and should) be used to adapt any binary classification metric to the multiclass classification task.

Evaluating OvO and OvR results also can help understanding which classes the model is struggling to describe, and which features you can add or remove to improve the result of the model.

#### **Cross Validation of Logistic Regression Model Using Stratified K\_fold method**

- Stratified kfold cross validation is an extension of regular kfold cross validation but specifically for classification problems where rather than the splits being completely random, the ratio between the target classes is the same in each fold as it is in the full dataset.
- Stratified kfold cross validation is typically useful when we have imbalanced data and where the data size is on the small side. Sometimes we will over or under sample our data to deal with class imbalance but other times we want to maintain the class imbalance when it's representative of or contains some information about what we are trying to classify.
- When the data is large enough we can still use regular kfold cross validation as this will often preserve the class ratios but this becomes less likely with fewer training examples.
- Roughly and succinctly speaking (depending of course on the test, training and validation scheme you employ), things like std. dev. can be used as a rudimentary measure of classifier stability. If you you are using stratified k-fold cross validation (depending on how it is implemented), then the mean *could* represent the mean of the aggregated correct classification results over k models (there are other ways to interpret it of course). The std. dev would then give an indication as to how these varied across the different models. If you are not interested in outright classification performance, but merely want to compare the stability of two different learners (again in a rather rudimentary fashion), you could use something like relative std. dev. = std dev/correct classification percentage.

In this project, I have choosen Stratified kfold cross validation technique, because the ratio between the target classes is the same in each fold as it is in the full dataset after performing the "RandomOversampling" technique.

```
In [53]: from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import cross_val_score
```

```
In [54]: skf = StratifiedKFold(n_splits = 10, shuffle = True, random_state = None)
```

```
In [55]: lr_cvscr = cross_val_score(model1, x, y, cv = skf)
```

```
In [56]: lr_cvscr
```

```
Out[56]: array([0.85915493, 0.87793427, 0.85915493, 0.87323944, 0.92957746,
0.89671362, 0.89150943, 0.85377358, 0.88679245, 0.86792453])
```

```
In [57]: # Check the Accuracy
lr_cvscr_Accuracy = lr_cvscr * 100.0
lr_cvscr_Accuracy
```

```
Out[57]: array([85.91549296, 87.79342723, 85.91549296, 87.32394366, 92.95774648,
89.6713615 , 89.1509434 , 85.37735849, 88.67924528, 86.79245283])
```

```
In [58]: lr_cvscr_Accuracy_mean = lr_cvscr_Accuracy.mean()          # Average of Crossvalidation score
```

```
In [59]: lr_cvscr_Accuracy_std = lr_cvscr_Accuracy.std()           # Standard deviation of crossvalidation score.
```

### ***KNeighbours Clssifier***

```
In [60]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [61]: model2 = KNeighborsClassifier(n_neighbors = 3)
```

```
In [62]: model2.fit(x_train, y_train)
```

```
Out[62]: 

KNeighborsClassifier



KNeighborsClassifier(n_neighbors=3)


```

```
In [63]: kn_y_pred = model2.predict(x_test)
```

```
In [64]: cm2 = confusion_matrix(y_test, kn_y_pred)
cm2.ravel()
```

```
Out[64]: array([442, 48, 6, 8, 485, 4, 0, 2, 495], dtype=int64)
```

```
In [65]: pred2 = model2.predict_proba(x_test)
logloss2 = log_loss(kn_y_pred, pred2)          #Logloss score
```

```
In [66]: acc2 = accuracy_score(y_test, kn_y_pred)
```

- When Parameter is set to n\_neighbors = 3 gives 95%, n\_neighbors = 5 gives 92%, n\_neighbors = 7 gives 91%. So, i kept first parameter for best accuracy score

```
In [67]: precision_positive2 = metrics.precision_score(y_test, kn_y_pred, pos_label = 'positive', average='micro')
precision_negative2 = metrics.precision_score(y_test, kn_y_pred, pos_label = 'negative', average='micro')
```

```
In [68]: recall_sensitivity2 = metrics.recall_score(y_test, kn_y_pred, pos_label = 'positive', average='micro') # Recall_Score
recall_specificity2 = metrics.recall_score(y_test, kn_y_pred, pos_label = 'negative', average='micro')
```

```
In [69]: f1_positive2 = metrics.f1_score(y_test, kn_y_pred, pos_label = 'positive', average='micro') # f1_Score
f1_negative2 = metrics.f1_score(y_test, kn_y_pred, pos_label = 'negative', average='micro')
```

```
In [70]: rocscore2 = roc_auc_score(y_test, pred2, multi_class='ovr', average= 'weighted')
```

### **Crossvalidation of KNeighborsClassifier Model**

```
In [71]: kn_cvscr = cross_val_score(model2, x, y, cv = skf)
```

```
In [72]: # Check the Accuracy
kn_cvscr_Accuracy = kn_cvscr * 100.0
kn_cvscr_Accuracy
```

```
Out[72]: array([94.83568075, 88.73239437, 92.01877934, 91.54929577, 90.14084507,
                92.48826291, 87.73584906, 89.1509434 , 88.20754717, 92.45283019])
```

```
In [73]: kn_cvscr_Accuracy_mean = kn_cvscr_Accuracy.mean() # Average of Crossvalidation score
```

```
In [74]: kn_cvscr_Accuracy_std = kn_cvscr_Accuracy.std() # Standard deviation of crossvalidation score.
```

### **SupportVectorMachine**

```
In [75]: from sklearn.svm import SVC
model3 = SVC(kernel = 'rbf', decision_function_shape='ovr', gamma = 'auto', probability = True)
```

```
In [76]: model3.fit(x_train, y_train)
```

```
Out[76]: SVC(gamma='auto', probability=True)
```

```
In [77]: svc_y_pred = model3.predict(x_test)
```

```
In [78]: cm3 = confusion_matrix(y_test, svc_y_pred)
cm3.ravel()
```

```
Out[78]: array([492,   3,   1,   0, 494,   3,   0,   0, 497], dtype=int64)
```

```
In [79]: acc3 = accuracy_score(y_test, svc_y_pred)
```

- 1 When SVC(kernel = 'rbf', decision\_function\_shape='ovr'), the accuracy score is 81%
- 2 When svc(kernel = 'linear', decision\_function\_shape='ovr'), the accuracy score is 84%
- 3 When svc(kernel = 'poly', decision\_function\_shape='ovr'), the accuracy score is 82%
- 4 When svc(kernel = 'sigmoid', decision\_function\_shape='ovr'), the accuracy score is 32% means very poor.
- 5 When svc(kernel = 'rbf', decision\_function\_shape='ovr', gamma='scale'), the accuracy score is 81%
- 6 When svc(kernel = 'rbf', decision\_function\_shape='ovr', gamma='auto'), the accuracy score is 98%
- 7 When svc(kernel = 'linear', decision\_function\_shape='ovr', gamma='scale'), the accuracy score is 84%
- 8 When svc(kernel = 'linear', decision\_function\_shape='ovr', gamma='auto'), the accuracy score is 84%
- 9 When svc(kernel = 'poly', decision\_function\_shape='ovr', gamma='scale'), the accuracy score is 82%
- 10 When svc(kernel = 'poly', decision\_function\_shape='ovr', gamma='auto'), the accuracy score is 90%
- 11 When svc(kernel = 'sigmoid', decision\_function\_shape='ovr', gamma='scale'), the accuracy score is 31%
- 12 When svc(kernel = 'sigmoid', decision\_function\_shape='ovr', gamma='auto'), the accuracy score is 33%

Among all above parameter tuning 6th one gives best accuracy score, For this model i selected 6th one.

```
In [80]: pred3 = model3.predict_proba(x_test)
logloss3 = log_loss(svc_y_pred, pred3)           #Logloss score
```

```
In [81]: precision_positive3 = metrics.precision_score(y_test, svc_y_pred, pos_label = 'positive', average='micro')
precision_negative3 = metrics.precision_score(y_test, svc_y_pred, pos_label = 'negative', average='micro')
```

```
In [82]: recall_sensitivity3 = metrics.recall_score(y_test, svc_y_pred, pos_label = 'positive', average='micro') # Recall_Score
recall_specificity3 = metrics.recall_score(y_test, svc_y_pred, pos_label = 'negative', average='micro')
```

```
In [83]: f1_positive3 = metrics.f1_score(y_test, svc_y_pred, pos_label = 'positive', average='micro')           # f1_Score
f1_negative3 = metrics.f1_score(y_test, svc_y_pred, pos_label = 'negative', average='micro')
```

```
In [84]: rocscore3 = roc_auc_score(y_test, pred3, multi_class='ovr', average= 'weighted')
```

### **Cross Validation of SVM**

```
In [85]: svc_cvscr = cross_val_score(model3, x, y, cv = skf)
```

```
In [86]: # Check the Accuracy
svc_cvscr_Accuracy = svc_cvscr * 100.0
svc_cvscr_Accuracy
```

```
Out[86]: array([79.342723 , 79.342723 , 81.22065728, 78.40375587, 79.342723 ,
               81.22065728, 82.0754717 , 80.66037736, 80.66037736, 79.24528302])
```

```
In [87]: svc_cvscr_Accuracy_mean = svc_cvscr_Accuracy.mean()           # Average of Crossvalidation score
```

```
In [88]: svc_cvscr_Accuracy_std = svc_cvscr_Accuracy.std()             # Standard deviation of crossvalidation score.
```

### **Decision Tree**

```
In [89]: from sklearn.tree import DecisionTreeClassifier
model4 = DecisionTreeClassifier(criterion = 'log_loss')
```

```
In [90]: model4.fit(x_train,y_train)
```

```
Out[90]: ▾      DecisionTreeClassifier
DecisionTreeClassifier(criterion='log_loss')
```

```
In [91]: dt_y_pred = model4.predict(x_test)
```

```
In [92]: cm4 = confusion_matrix(y_test, dt_y_pred)
cm4.ravel()
```

```
Out[92]: array([472,  18,   6,   0, 497,   0,   0,   0, 497], dtype=int64)
```

```
In [93]: acc4 = accuracy_score(y_test, dt_y_pred)
```

- When DecisionTreeClassifier(criterion = 'entropy'), the accuracy is 98%
- When DecisionTreeClassifier(criterion = 'gini'), the accuracy is 97%
- When DecisionTreeClassifier(criterion = 'log\_loss'), the accuracy is 98%

Among all above three, I've Chosen last one.

```
In [94]: pred4 = model4.predict_proba(x_test)
logloss4 = log_loss(dt_y_pred, pred4)           #Logloss score
```

```
In [95]: precision_positive4 = metrics.precision_score(y_test, dt_y_pred, pos_label = 'positive', average='micro')
precision_negative4 = metrics.precision_score(y_test, dt_y_pred, pos_label = 'negative', average='micro')
```

```
In [96]: recall_sensitivity4 = metrics.recall_score(y_test, dt_y_pred, pos_label = 'positive', average='micro') # Recall_Score

recall_specificity4 = metrics.recall_score(y_test, dt_y_pred, pos_label = 'negative', average='micro')
```

```
In [97]: f1_positive4 = metrics.f1_score(y_test, dt_y_pred, pos_label = 'positive', average='micro') # f1_Score

f1_negative4 = metrics.f1_score(y_test, dt_y_pred, pos_label = 'negative', average='micro')
```

```
In [98]: rocscore4 = roc_auc_score(y_test, pred4, multi_class='ovr', average= 'weighted')
```

### **Crossvalidation of DecisionTreeClasssifier**

```
In [99]: dt_cvscr = cross_val_score(model4, x, y, cv = skf)
```

```
In [100]: # Check the Accuracy
dt_cvscr_Accuracy = dt_cvscr * 100.0
dt_cvscr_Accuracy
```

```
Out[100]: array([90.14084507, 92.95774648, 93.42723005, 93.89671362, 91.07981221,
                95.77464789, 94.81132075, 94.33962264, 93.86792453, 92.45283019])
```

```
In [101]: dt_cvscr_Accuracy_mean = dt_cvscr_Accuracy.mean() # Average of Crossvalidation score
```

```
In [102]: dt_cvscr_Accuracy_std = dt_cvscr_Accuracy.std() # Standard deviation of crossvalidation score.
```

### **NaiveBayes**

```
In [103]: from sklearn.naive_bayes import GaussianNB
model5 = GaussianNB()
```

```
In [104]: model5.fit(x_train, y_train)
```

```
Out[104]: ▾ GaussianNB
GaussianNB()
```

```
In [105]: nb_y_pred = model5.predict(x_test)
```

```
In [106]: cm5 = confusion_matrix(y_test, nb_y_pred)
cm5.ravel()
```

```
Out[106]: array([362, 113, 21, 30, 435, 32, 21, 142, 334], dtype=int64)
```

```
In [107]: acc5 = accuracy_score(y_test, nb_y_pred)
```

```
In [108]: pred5 = model5.predict_proba(x_test)
logloss5 = log_loss(nb_y_pred, pred5) #Logloss score
```

```
In [109]: precision_positive5 = metrics.precision_score(y_test, nb_y_pred, pos_label = 'positive', average='micro')

precision_negative5 = metrics.precision_score(y_test, nb_y_pred, pos_label = 'negative', average='micro')

recall_sensitivity5 = metrics.recall_score(y_test, nb_y_pred, pos_label = 'positive', average='micro') # Recall_Score

recall_specificity5 = metrics.recall_score(y_test, nb_y_pred, pos_label = 'negative', average='micro')

f1_positive5 = metrics.f1_score(y_test, nb_y_pred, pos_label = 'positive', average='micro') # f1_Score

f1_negative5 = metrics.f1_score(y_test, nb_y_pred, pos_label = 'negative', average='micro')

rocscore5 = roc_auc_score(y_test, pred5, multi_class='ovr', average= 'weighted')
```

### **Crossvalidation of Naive Bayes**

```
In [110]: nb_cvscr = cross_val_score(model5, x, y, cv = skf)
```

```
In [111]: # Check the Accuracy
nb_cvscr_Accuracy = nb_cvscr * 100.0
nb_cvscr_Accuracy
```

```
Out[111]: array([80.75117371, 82.62910798, 82.62910798, 80.28169014, 83.09859155,
              76.99530516, 80.66037736, 82.0754717 , 78.30188679, 85.37735849])
```

```
In [112]: nb_cvscr_Accuracy_mean = nb_cvscr_Accuracy.mean() # Average of Crossvalidation score
nb_cvscr_Accuracy_std = nb_cvscr_Accuracy.std() # Standard deviation of crossvalidation score.
```

### **Randomforesrt**

```
In [113]: from sklearn.ensemble import RandomForestClassifier
model6 = RandomForestClassifier(n_estimators = 100, criterion = 'log_loss')
```

```
In [114]: model6.fit(x_train, y_train)
```

```
Out[114]: ▼      RandomForestClassifier
RandomForestClassifier(criterion='log_loss')
```

```
In [115]: rf_y_pred = model6.predict(x_test)
```

```
In [116]: cm6 = confusion_matrix(y_test, rf_y_pred)
cm6.ravel()
```

```
Out[116]: array([478, 13, 5, 0, 497, 0, 0, 0, 497], dtype=int64)
```

```
In [117]: acc6 = accuracy_score(y_test, rf_y_pred)
```

- When RandomForestClassifier(n\_estimators = 100, criterion = 'gini'), the accuracy is 98.5%

- When RandomForestClassifier(n\_estimators = 100, criterion = 'entropy'), the accuracy is 98.7%
- When RandomForestClassifier(n\_estimators = 100, criterion = 'log\_loss'), the accuracy is 98.6%

Among all of the above, I have choosen second one.

```
In [118]: pred6 = model6.predict_proba(x_test)
logloss6 = log_loss(rf_y_pred, pred6)           #logloss score
```

```
In [119]: precision_positive6 = metrics.precision_score(y_test, rf_y_pred, pos_label = 'positive', average='micro')
precision_negative6 = metrics.precision_score(y_test, rf_y_pred, pos_label = 'negative', average='micro')
recall_sensitivity6 = metrics.recall_score(y_test, rf_y_pred, pos_label = 'positive', average='micro') # Recall_Score
recall_specificity6 = metrics.recall_score(y_test, rf_y_pred, pos_label = 'negative', average='micro')
f1_positive6 = metrics.f1_score(y_test, rf_y_pred, pos_label = 'positive', average='micro')           # f1_Score
f1_negative6 = metrics.f1_score(y_test, rf_y_pred, pos_label = 'negative', average='micro')
rocscore6 = roc_auc_score(y_test, pred6, multi_class='ovr', average= 'weighted')
```

### **Crossvalidation of Randomforest**

```
In [120]: rf_cvscr = cross_val_score(model6, x, y, cv = skf)
```

```
In [121]: # Check the Accuracy
rf_cvscr_Accuracy = kn_cvscr * 100.0
rf_cvscr_Accuracy
```

```
Out[121]: array([94.83568075, 88.73239437, 92.01877934, 91.54929577, 90.14084507,
                92.48826291, 87.73584906, 89.1509434 , 88.20754717, 92.45283019])
```

```
In [122]: rf_cvscr_Accuracy_mean = rf_cvscr_Accuracy.mean()           # Average of Crossvalidation score
rf_cvscr_Accuracy_std = rf_cvscr_Accuracy.std()           # Standard deviation of crossvalidation score.
```

### **Display of All scores**



```
In [123]: data = {"Accuracy_Score": [acc1, acc2, acc3, acc4, acc5, acc6],
                  "log_loss_Score": [logloss1, logloss2, logloss3, logloss4, logloss5, logloss6],
                  "F1_Score_Positive": [f1_positive1, f1_positive2, f1_positive3, f1_positive4, f1_positive5, f1_positive6],
                  "F1_Score_Negative": [f1_negative1, f1_negative2, f1_negative3, f1_negative4, f1_negative5, f1_negative6],
                  "Precision_Positive": [precision_positive1, precision_positive2, precision_positive3, precision_positive4, precision_positive5, precision_positive6],
                  "Precision_Negative": [precision_negative1, precision_negative2, precision_negative3, precision_negative4, precision_negative5, precision_negative6],
                  "Recall_sensitivity": [recall_sensitivity1, recall_sensitivity2, recall_sensitivity3, recall_sensitivity4, recall_sensitivity5, recall_sensitivity6],
                  "Recall_specificity": [recall_specificity1, recall_specificity2, recall_specificity3, recall_specificity4, recall_specificity5, recall_specificity6],
                  "ROC_Score": [rocscore1, rocscore2, rocscore3, rocscore4, rocscore5, rocscore6],
                  "Cross_Val Score": [lr_cvscr_Accuracy_mean, kn_cvscr_Accuracy_mean, svc_cvscr_Accuracy_mean, dt_cvscr_Accuracy_mean, nb_cvscr_Accuracy_mean, rf_cvscr_Accuracy_mean],
                  "Cross_Val Std": [lr_cvscr_Accuracy_std, kn_cvscr_Accuracy_std, svc_cvscr_Accuracy_std, dt_cvscr_Accuracy_std, nb_cvscr_Accuracy_std, rf_cvscr_Accuracy_std]}
All_scores = pd.DataFrame(data, index = ["LR", "KNN", "SVM", "DT", "NB", "RF"])
```

```
In [124]: All_scores
```

```
Out[124]:
```

	Accuracy_Score	log_loss_Score	F1_Score_Positive	F1_Score_Negative	Precision_Positive	Precision_Negative	Recall_sensitivity	Recall_specificity	ROC_Score	Cross_Val Score	Cross_Val Std
<b>LR</b>	0.821477	4.445686e-01	0.821477	0.821477	0.821477	0.821477	0.821477	0.821477	0.943090	87.957746	2.160263
<b>KNN</b>	0.954362	2.013719e-02	0.954362	0.954362	0.954362	0.954362	0.954362	0.954362	0.981083	90.731243	2.177317
<b>SVM</b>	0.995302	1.276138e-02	0.995302	0.995302	0.995302	0.995302	0.995302	0.995302	0.998256	80.151475	1.111034
<b>DT</b>	0.983893	2.109424e-15	0.983893	0.983893	0.983893	0.983893	0.983893	0.983893	0.987915	93.274869	1.610018
<b>NB</b>	0.759060	5.232484e-02	0.759060	0.759060	0.759060	0.759060	0.759060	0.759060	0.887387	81.280007	2.307738
<b>RF</b>	0.987919	3.457891e-02	0.987919	0.987919	0.987919	0.987919	0.987919	0.987919	0.999538	90.731243	2.177317

### Conclusion

- In this project i have applied Classification algorithms to find out best model to prdict output.
- Among all the models i have selected SVM on the basis of Accuracy score and Cross Validation score.

```
In [125]: joblib.dump(model3, "regressor.pkl")
```

```
Out[125]: ['regressor.pkl']
```

```
In [126]: check = fetal.copy()
```

In [127]:

```
check
```

Out[127]:

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations	prolongued_decelerations	abnormal_short_term_variability	mean_value_of_short_term_variability	percentage_of_time_with
0	120	0.000	0.000	0.000	0.000	0.0	0.0	73	0.5	
1	132	0.006	0.000	0.006	0.003	0.0	0.0	17	2.1	
2	133	0.003	0.000	0.008	0.003	0.0	0.0	16	2.1	
3	134	0.003	0.000	0.008	0.003	0.0	0.0	16	2.4	
4	132	0.007	0.000	0.008	0.000	0.0	0.0	16	2.4	
...	...	...	...	...	...	...	...	...	...	
2121	140	0.000	0.000	0.007	0.000	0.0	0.0	79	0.2	
2122	140	0.001	0.000	0.007	0.000	0.0	0.0	78	0.4	
2123	140	0.001	0.000	0.007	0.000	0.0	0.0	79	0.4	
2124	140	0.001	0.000	0.006	0.000	0.0	0.0	78	0.4	
2125	142	0.002	0.002	0.008	0.000	0.0	0.0	74	0.4	

2126 rows × 22 columns

In [128]: `check.drop("fetal_health",axis=1,inplace=True)`

In [129]:

```
check
```

Out[129]:

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations	prolongued_decelerations	abnormal_short_term_variability	mean_value_of_short_term_variability	percentage_of_time_with
0	120	0.000	0.000	0.000	0.000	0.0	0.0	73	0.5	
1	132	0.006	0.000	0.006	0.003	0.0	0.0	17	2.1	
2	133	0.003	0.000	0.008	0.003	0.0	0.0	16	2.1	
3	134	0.003	0.000	0.008	0.003	0.0	0.0	16	2.4	
4	132	0.007	0.000	0.008	0.000	0.0	0.0	16	2.4	
...	...	...	...	...	...	...	...	...	...	
2121	140	0.000	0.000	0.007	0.000	0.0	0.0	79	0.2	
2122	140	0.001	0.000	0.007	0.000	0.0	0.0	78	0.4	
2123	140	0.001	0.000	0.007	0.000	0.0	0.0	79	0.4	
2124	140	0.001	0.000	0.006	0.000	0.0	0.0	78	0.4	
2125	142	0.002	0.002	0.008	0.000	0.0	0.0	74	0.4	

2126 rows × 21 columns

In [130]: `final = check.iloc[:,:]`

In [131]:

```
final
```

Out[131]:

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations	prolongued_decelerations	abnormal_short_term_variability	mean_value_of_short_term_variability	percentage_of_time_with_ab
0	120	0.0	0.0	0.0	0.0	0.0	0.0	73	0.5	

1 rows × 21 columns

In [132]:

```
final.columns
```

Out[132]: Index(['baseline value', 'accelerations', 'fetal\_movement',  
'uterine\_contractions', 'light\_decelerations', 'severe\_decelerations',  
'prolongued\_decelerations', 'abnormal\_short\_term\_variability',  
'mean\_value\_of\_short\_term\_variability',  
'percentage\_of\_time\_with\_abnormal\_long\_term\_variability',  
'mean\_value\_of\_long\_term\_variability', 'histogram\_width',  
'histogram\_min', 'histogram\_max', 'histogram\_number\_of\_peaks',  
'histogram\_number\_of\_zeroes', 'histogram\_mode', 'histogram\_mean',  
'histogram\_median', 'histogram\_variance', 'histogram\_tendency'],  
dtype='object')

In [133]:

```
final_model = joblib.load("regressor.pkl")  
final_model
```

Out[133]:

▼

SVC

SVC(gamma='auto', probability=True)

In [134]:

```
final_model.predict(final)
```

Out[134]: array([2], dtype=int64)

In [ ]: